

Strategies for Collecting, Processing, and Analyzing Tweets from Large Newsworthy Events

Nick Ruest
York University
Toronto
ruestn@yorku.ca

KEYWORDS

Web archiving, Twitter

ACM Reference format:

Nick Ruest. 1997. Strategies for Collecting, Processing, and Analyzing Tweets from Large Newsworthy Events. In *Proceedings of Web Archiving and Digital Libraries 2017, Toronto, ON, June 2017 (JCDL 2017)*, 1 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

#WomensMarch, #Aleppo, #paris, #bataclan, #parisattacks, #porteouverte, #jesuischarlie, #jesuisahmed, #jesuisjuif, #charliehebd, #panamanpapers, and #eln42 are all different hashtags, but they share several things in common. They are all large newsworthy events. They are datasets that each contain over a million tweets. Most importantly these collections raise some interesting insights in collecting, processing, and analyzing large newsworthy events[2].

Collecting tweets from these events can be challenging because of timing. Tweets can be collected from the Filter API¹ and Search API². Both having their own caveats. The Filter API only captures the current Twitter stream, and is limited to collecting up to 1% of the overall Twitter stream. The Search API allows you to collect more than 1% of the overall Twitter stream[1], but one can only collect up to 18,000 every 15 minutes, and is limited to a 7 day window. Generally, using a strategy of using the Filter and Search API to capture a given event is the best.

DocNow's twarc³ includes a number of utilities to process a dataset after collection. These tools allow a researcher, librarian, or archivist to filter their dataset(s) down to what is needed for appraisal, and then accession. Noteworthy tools include; deduplication, source, retweets, date/times, users, and hashtags.

DocNow's utilities⁴ can be further used to curate related collections. One can extract all the urls of a dataset, unshorten them, and extract the unique urls to use as a seed list for a web crawler to capture websites related to a given event. One can also extract all of

the image urls, and download all images associated with a dataset, which then can be used for image analysis[3], presentation, and/or preservation.

In conclusion, this presentation will provide an overview of collection strategy, insights from processing and analysis, ensuing web crawls, and image presentation from each collection.

REFERENCES

- [1] Kevin Driscoll and Shawn Walker. 2014. Working Within a Black Box: Transparency in the Collection and Production of Big Twitter Data. *International Journal of Communication* 8 (2014), 1745–1764.
- [2] Ian Milligan, Nick Ruest, and Jimmy Lin. 2016. Content Selection and Curation for Web Archiving: The Gatekeepers vs. The Masses. In *Proceedings of the 16th ACM/IEEE-CS on Joint Conference on Digital Libraries (JCDL '16)*. ACM, New York, NY, USA, 107–110. <https://doi.org/10.1145/2910896.2910913>
- [3] Nick Ruest. 2016. 1,203,867 #eln42 images. (March 2016). <http://ruebot.net/post/1203867-eln42-images>

¹<https://dev.twitter.com/streaming/public>

²<https://dev.twitter.com/rest/public/search>

³<http://github.com/docnow/twarc>

⁴<https://github.com/DocNow/twarc/tree/master/utlis>

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

JCDL 2017, June 2017, Toronto, ON

© 2017 Copyright held by the owner/author(s).

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM.

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>